

Tests as Classifiers

Moshe Zadka – <https://cobordism.com>

Acknowledgement of Country

Belmont (in San Francisco Bay Area Peninsula)
Ancestral homeland of the Ramaytush Ohlone

What is a classifier?

Input: Something

What is a classifier?

Input: Something

Output: True/False

What is a classifier?

Input: Something

Output: True/False

Belongs to set / Does not belong to set

Classifier example: chihuahua or muffin

Input: Image of chihuahua or muffin

Classifier example: chihuahua or muffin

Input: Image of chihuahua or muffin

Output: Is it a dog?

Classifier failure: False alarm

Classifier says "yes", should say "no"

Example: Picture of muffin, classifier says "dog"

Classifier failure: Missing alarm

Classifier says "no", should say "yes"

Example: Picture of dog, classifier says "muffin"

Test (suites) as classifiers

Input: Code change

Output: Is the code buggy?

Test (suites) as classifiers

Tests suite failure: "Code buggy"

Test suite success: "Code not buggy"

Simple classifiers

Always alarm: "Yes" regardless of input

Never alarm: "No" regardless of input

Always alarm: a test

```
def test_always_alarm():  
    assert 1 == 0
```

Never alarm: a test suite

Empty file

Why not simple classifiers?

Writing tests is hard work!

Why not simple classifiers?

Writing tests is hard work!
Can we quantify the value?

Precision

Rewards not alarming:

```
precision = (  
    true_alarms /  
    (true_alarms + false_alarms)  
)
```

Recall

Rewards alarming:

```
recall = (  
    true_alarms /  
    (true_alarms + missing_alarms)  
)
```

Balancing Precision and Recall: F score

Harmonic mean:

$$\text{precision_inv} = 1 / \text{precision}$$

$$\text{recall_inv} = 1 / \text{recall}$$

$$\text{mean_inv} = (\text{precision_inv} + \text{recall_inv}) / 2$$

$$\text{f_score} = 1 / \text{mean_inv}$$

Balancing Precision and Recall: F beta score

What if it's not equally important?

Balancing Precision and Recall: F beta score

What if it's not equally important?

```
precision_inv = 1 / precision
recall_inv = 1 / recall
mean_inv = (
    (beta ** 2 * precision_inv + recall_inv)
    /
    (beta ** 2 + 1)
f_beta_score = 1 / mean_inv
```

Balancing Precision and Recall: Who is beta?

F score is F beta score when beta is 1

Balancing Precision and Recall: Who is beta?

F score is F beta score when beta is 1

The "beta" parameter encodes utility: which error hurts harder?

Balancing Precision and Recall: Who is beta?

F score is F beta score when beta is 1

The "beta" parameter encodes utility: which error hurts harder?
A business decision!

Beta: a Meaning

Beta is 2: Missing alarms are twice as painful as false alarms

Beta: a Meaning

Beta is 2: Missing alarms are twice as painful as false alarms

Beta is 0.5: Missing alarms are half as painful as false alarms

Tests and the F score

What makes a test bring down the F score?

False Alarm: Flakey test

Fails "randomly"

False Alarm: Implementation test

Testing implementation details:

False Alarm: Implementation test

Testing implementation details:

```
def test_implementation():  
    with mock.patch(  
        "subprocess.check_output"  
    ) as check_output:  
        run_code()  
        assert some_stuff
```

What happens when it uses "subprocess.run"?

Missing Alarm: Non-covered code

What is not run does not affect result of tests

Missing Alarm: Loose assertions

Asserting something is greater than 5,

Missing Alarm: Loose assertions

Asserting something is greater than 5,
not equal to 6

Estimating F score: Flakey tests

Run known-good main branch

Estimating F score: Flakey tests

Run known-good main branch
check for failures

Estimating F score: Implementation tests

Rough measure: changes to tests

Estimating F score: Implementation tests

Rough measure: changes to tests

Heuristics to compensate for legitimate changes

Estimating F score: Mutation testing

Percentage of surviving mutants: missing alarms

Estimating F score: Check reported bugs

Bugs with added tests

Estimating F score: Check reported bugs

Bugs with added tests
but not features

Estimating F score: Check reported bugs

Bugs with added tests
but not features
Heuristics

Estimating F score: True alarms

Weighing by PR/branch

Estimating F score: True alarms

Weighing by PR/branch

Try and get data from dev machines too!

Estimating F score: Lagging indicator

Merged PRs / main development

F score cautions

Check heuristics

F score cautions

Check heuristics

Goodhart's law

F score usage

Revealed preferences: beta

F score usage

Revealed preferences: beta

Calibrate time investment in improving tests

Improving tests

F score: lagging guide

Improve: Flakey tests

Better isolation

Improve: Implementation tests

Clear contracts

Improve: Untested code

Coverage

Improve: Under-tested code

Mutation testing

Summary: Goal

Prevent bugs

Summary: Goal

Prevent bugs
minimal cost

Summary: Decide

Cost of bug

Summary: Measure

Expectations vs. Reality

Summary: Improve

Align reality